

M&A OPERATING SYSTEM

Product Design

AI Analytics Platform

Draft v1.0 · May 2026

Andrew Bush (www.maoperatingsystem.com/bio-andrew-bush)

1. Core Platform Capabilities

This chapter is the logical architecture reference for the AI Analytics Platform. It covers nine pipeline components in the order a query encounters them, using a single portfolio manager query as a running example throughout. Each section describes what a component does, its controls contract, and where in the pipeline it sits — with no references to specific technology products or vendor implementations.

The pipeline processes every request in this sequence: the **MCP Capability Layer** is the entry point — the governed API through which all consumers access the platform. The **Semantic Metrics Repository** is the governing catalogue — every queryable metric, dimension, and operation must be registered here before it is resolvable. The **Intent Resolution Agent** receives natural language queries, retrieves candidate operations from the SMR via RAG, and uses a language model to rank and resolve the user's intent into a structured operation and parameters. The **Role-Aware Projection Layer** receives the resolved request and the caller's JWT, makes the entitlement decisions, and produces an entitlement projection before any query plan is compiled. The **Semantic Validation Layer** validates the resolved structured request, enforces the entitlement projection, and produces a platform-agnostic query plan. The **Semantic Controls Layer** applies thresholds and compliance classification before releasing the plan for execution. The **Physical Query Planner** translates the approved logical plan into backend-specific physical query fragments. The **Federated Query Engine** routes those fragments to registered backends, executes in parallel, and assembles the result. The **Data Visualization Language (DVL)** selects the presentation contract. The **Narrative Synthesis Agent** produces a governed plain-language summary via a tightly-scoped language model call. The **Analytical Lineage Store** records the complete computation provenance. The **Provenance Artifact Service** assembles and seals the compliance artifact when a query meets the two-signal compliance trigger.

Platform roles — who interacts with each component and how — are defined before the component descriptions.

Platform Roles

The platform operates across three distinct planes: an **analytical plane** (querying, exploring, and exporting governed data), a **controls plane** (defining, approving, and administering the semantic layer and its access controls), and an **infrastructure plane** (platform deployment, health, and technical configuration).

Role	Plane	Definition
Analytical End User	Analytical	Ask governed analytical questions via natural language; receive role-constrained results without knowledge of data structures or metric identifiers
Power Analyst	Analytical	Multi-dimensional exploration, governed drilldown, lineage inspection, result export

Role	Plane	Definition
Data Modeller	Controls	Owns semantic data definitions in the SDR: logical data elements, object models, business definitions, critical data elements, and physical schema mappings. Ensures the organisation's data assets are accurately described and structured — the foundational layer on which metric definitions are built
Metrics Modeller	Controls	Owns semantic metrics and analytics definitions in the SMR: key performance metrics, analytics operations, trend analysis constructs, and insight definitions. Must combine domain knowledge — what does this metric mean in this business context — with modelling precision: how it is calculated, from which sources, under which dimensional hierarchies, and with which access policies
Entitlements Manager	Controls	Responsible for defining and maintaining the organisation's data entitlement policies: who may perform which actions (create, read, update, delete) on which data elements, analytics definitions, and business process metrics. Configures the metric access sets, dimension access sets, row predicates, and column masks that RAPL enforces at query time
Analytics Governance	Controls	Overall accountability for the governance, integrity, and outcomes of the analytics platform. Owns SMR registry health, approves semantic definition changes from Metrics Modellers, oversees entitlement policy governance, and is accountable for the quality, accuracy, and completeness of analytical outputs across the organisation. Reviews platform success metrics and controls health indicators. The final authority on what is defined, who can access it, and whether the platform is delivering the right outcomes. Must be in place before go-live — without this role, the registry has no approval authority and the platform cannot serve any query
Integration Engineer	Controls	Registers execution backends, maintains connection configuration, and declares the physical mappings that the Federated Query Engine resolves at execution time. Operates through configuration interfaces only — not the query path
Platform Admin	Infrastructure	Infrastructure and operations team. Responsible for platform health, deployment, infrastructure-level governance, and technical platform configuration including controls settings, feature flags, and deployment configuration. Implements the technical policies and settings determined by Analytics Governance. Has no query interface into analytical data

Roles are not mutually exclusive. A single individual may hold multiple roles; the platform evaluates entitlements from the combined JWT claims present at query time.

The **Data Modeller** and **Metrics Modeller** are the critical pre-conditions for everything downstream. No analytical query can be served against a metric that has not been modelled, registered, and approved. The Data Modeller establishes the foundational data definitions in the SDR — without accurate data structure definitions, metric definitions cannot be built. The Metrics Modeller builds on that foundation to define the analytical layer in the SMR — without registered, approved metric definitions, the controls pipeline, the entitlement layer, the lineage store, and the Data Visualization Language (DVL) have nothing to operate on. Analytics Governance holds final approval authority over both layers.

Role × Feature Access


Feature	End User	Power Analyst	Data Modeller	Metrics Modeller	Entitlements Mgr	Analytics Governance	Integration Eng	Platform Admin
Natural language query	✓	✓		✓		✓		
Role-aware results	✓	✓		✓		✓		
Governed drilldown		✓		✓		✓		
Lineage inspector		✓	✓	✓		✓		
Result export		✓		✓		✓		
Provenance artifacts	✓	✓		✓		✓		
SMR browsing		✓	✓	✓	✓	✓		
SDR management			✓					
Metric definition authoring				✓				
Metric definition approval						✓		
Dimension & hierarchy management				✓		✓		
Entitlement policy management					✓	✓		
Backend registration							✓	
Controls configuration								✓

Feature	End User	Power Analyst	Data Modeller	Metrics Modeller	Entitlements Mgr	Analytics Governance	Integration Eng	Platform Admin
Audit trail					✓	✓	✓	✓
Platform infrastructure								✓

Architecture and Request Flow

The platform exposes its capability through three consumption modes. The first is direct API access, where a host-built custom analytics UI calls the MCP Capability Layer directly with a structured tool invocation, supplying a JWT for entitlement resolution and receiving a structured response containing a display specification and narrative. The second is conversational backend access, where the AI Chat Platform's conversation engine calls the Analytics Platform as a tool provider, mediating between a conversational UI component and the controls pipeline. The third is agentic access, where scheduled agents, event monitors, and automated report pipelines call the MCP Capability Layer with machine-issued JWTs to perform periodic or event-driven analytical tasks without human-in-the-loop interaction. These three modes share a single entry point and a single controls pipeline; the consumption mode affects only the caller's interaction pattern, not the trust model applied.


Architecture Diagram



The Analytics Engine is a single MCP server. It exposes three analytical tools (`run_analytics` , `list_operations` , `drilldown`) through a single MCP Capability Layer endpoint. The engine contains exactly two bounded AI steps: the Intent Resolution Agent (IRA), which identifies the right governed operation from the user's natural language query, and the Narrative Synthesis Agent (NSA), which summarises the computed result in plain text after execution. All stages between them — SVL, RAPL, SCL, PQP, and FQE — are entirely

For conversational consumers, the user's natural language query is forwarded directly to the Analytics Engine. Intent resolution — selecting the right governed operation and binding parameters — happens inside the engine's Intent Resolution Agent. The Analytics Engine returns the display specification, structured data, and governed narrative; the AI Chat Platform renders the result. Structured API consumers (agents, custom UIs) may call `run_analytics` with an explicit `operation_id` and `params`, bypassing the IRA entirely.

Request Flow



Running example: This query is traced through every component below — each section's Example shows the same request at the next stage of the pipeline.

```
"Show me portfolio returns versus benchmark for my equity portfolios this quarter."
```

AI Consumers

The Analytics Engine is accessed by three consumer types: a conversational AI platform (which mediates between a user and the governed query pipeline), autonomous agents and pipelines (which call the MCP layer directly with structured requests), and custom applications (which call the MCP layer with host-issued tokens). The AI Chat Platform is the conversational layer through which users interact with the Analytics Engine. Its role is to relay the user's natural language query to the Analytics Engine and render the structured result it receives back. Intent resolution — identifying which governed operation and parameters match the user's question — is performed inside the Analytics Engine by the Intent Resolution Agent, not by the consumer.

Natural language path. When a user asks an analytical question, the consumer forwards the natural language query and the user's JWT to the Analytics Engine. The engine's IRA handles operation selection, parameter binding, and — if intent is ambiguous — returns a confirmation card before proceeding to execution. The consumer does not need to know the SMR catalogue or construct structured parameters.

Structured path. Consumers that construct explicit `operation_id` + `params` payloads (agentic pipelines, custom analytics UIs, integration tests) call `run_analytics` with structured arguments directly. The `list_operations` tool returns the entitled operation catalogue for consumers that build their own operation selection UI. Structured calls bypass the IRA and route directly to the RAPL, which makes the entitlement decisions before the SVL compiles the plan.

Analytics Engine call. The tool call is submitted to the Analytics Engine with the user's JWT forwarded unmodified. The Analytics Engine validates, plans, executes, and returns a structured result and DVL display specification. Entitlement enforcement, query planning, and execution are entirely the Analytics Engine's responsibility.

Response assembly. The conversation engine renders the DVL display specification inline. When narrative synthesis is enabled, the `narrative` object in the response contains the governed summary produced by the Analytics Engine's NSA. The conversation engine surfaces this directly. The `result_id` is retained for any follow-up `drilldown` call or lineage inspection.

Example

```
"Show me portfolio returns versus benchmark for my equity portfolios this quarter."
```

↳ **Step 0 — Natural language query forwarded.** The user's question is relayed by the conversational AI directly to the Analytics Engine with the user's JWT. The consumer does not interpret, translate, or structure the query — it forwards it as-is.


```

POST /v1/mcp
Authorization: Bearer <host-issued-jwt>

{
  "jsonrpc": "2.0",
  "id": "req-8a3f2c",
  "method": "tools/call",
  "params": {
    "name": "run_analytics",
    "arguments": {
      "query": "Show me portfolio returns versus benchmark for my equity portfolios this quarter."
    }
  }
}

```

The Analytics Engine processes the request end-to-end and returns a structured result. The conversation engine renders the grouped bar chart from the DVL display specification and surfaces the governed narrative as the assistant's reply.

MCP Capability Layer (MCP)

Governing principles: P2 — Controls before execution · P5 — Role-aware by default

The MCP Capability Layer exposes the platform's governed analytical operations to AI orchestrators via MCP Streamable HTTP transport. Each capability is a bounded, named operation with a typed input schema, a governed execution path — at minimum through role-aware projection and the federated query engine, and through the full pipeline (RAPL → SVL → SCL → FQE) for metric and analytical operations, and a typed output contract. AI agents interact with capabilities, not databases. There is no privileged API path — AI agents receive the same controls-validated results as human users.

Tool Catalogue

The Analytics Engine exposes three tools. All analytical operations are SMR-catalogue driven — the code is the execution engine, not the operation registry. The SMR owns every operation definition: what parameters it needs, what metrics and dimensions it supports, and how deeply it runs through the pipeline via its `execution_profile`.

`run_analytics(operation_id: str, params: dict, jwt: str)` — Executes any SMR-registered operation. The operation's `execution_profile` in the SMR determines which pipeline stages run.

`list_operations(domain: str | None, jwt: str)` — Returns the SMR operation catalogue with operation IDs, display names, required parameters, supported metrics/dimensions, and execution profiles. Only operations the authenticated user is entitled to execute are returned.

`drilldown(result_id: str, hierarchy: str, selected_value: str | None, jwt: str)` — Navigates into a dimension hierarchy from a prior result. All filters, role predicates, and entitlement context from the original result are preserved.

Execution Profiles

Each SMR operation carries an `execution_profile` defined in its `analytical_operation` entry in the SMR catalogue. This tells the pipeline executor which stages to invoke. No execution depth is hardcoded in the MCP layer — it is always determined by the SMR catalogue.

Profile	Pipeline stages
<code>data_retrieval</code>	Auth → IRA → RAPL → FQE → Lineage
<code>metric_query</code>	Auth → IRA → RAPL → SVL → SCL → FQE → Lineage
<code>full_analytical</code>	Auth → IRA → RAPL → SVL → SCL → PQP → FQE → DVL + NSA + PAS → Lineage

Intent Confirmation Cards

When intent is ambiguous, or when `requiresIntentConfirmation: true` is set on the operation, the IRA returns candidate cards before executing any query. The cards are returned as the MCP response body in place of the analytical result. The consumer renders all candidates simultaneously; the user selects one, optionally refines it through the chat experience, then confirms to proceed.

Each card includes the resolved operation and parameters alongside a **presentation preview** — the anticipated chart type and axis structure — so the user can verify both what will be queried and how the result will be presented before execution commits.

The response payload uses a `candidates[]` array. A single-candidate response (governance override) uses the same structure with one entry:

```
{
  "intent_session_id": "ira-sess-20260605-wk4n",
  "candidates": [
    {
      "rank": 1,
      "confidence": 0.91,
      "operation_id": "compare_portfolios",
      "operation_label": "Compare Portfolios",
      "resolved_metrics": ["portfolio_return", "benchmark_return"],
      "resolved_dimensions": ["portfolio_id", "asset_class"],
      "time_period": "quarter_to_date",
      "filters": [{ "field": "asset_class", "operator": "eq", "value": "EQUITY" }],
      "estimated_performance_impact": 620,
      "classification": "INTERNAL",
      "presentation_hint": {
        "chart_type": "bar",
        "primary_dimension": "portfolio_id",
        "measures": ["portfolio_return", "benchmark_return"],
        "series_by": "metric"
      }
    },
    {
      "rank": 2,
      "confidence": 0.67,
      "operation_id": "portfolio_summary",
      "operation_label": "Portfolio Summary",
      "resolved_metrics": ["portfolio_return", "portfolio_nav"],
      "resolved_dimensions": ["portfolio_id"],
      "time_period": "quarter_to_date",
      "filters": [],
      "estimated_performance_impact": 280,
      "classification": "INTERNAL",
      "presentation_hint": {
        "chart_type": "table",
        "primary_dimension": "portfolio_id",
        "measures": ["portfolio_return", "portfolio_nav"],
        "series_by": null
      }
    }
  ],
  "action": "select or refine – re-submit with selected_candidate: <index> or a refinement message"
}
```

Selecting a candidate: re-submit `run_analytics` with `"selected_candidate": 0` (0-based index) to confirm the chosen interpretation and proceed to execution.

Refining through chat: send a natural language adjustment ("make it monthly, not quarterly") — the IRA updates the leading candidate's parameters and returns a revised card set. Refinement turns are bounded by `intentRefinementMaxTurns` (default 5) and recorded in the lineage record's `intent_session` field.

This is appropriate for any query where silent intent misresolution is unacceptable — including high-stakes, compliance-sensitive, or inherently ambiguous requests.

Capability Governance

Every capability invocation passes through the full controls pipeline: input schema validation → capability availability check (feature flags and role entitlements) → Role-Aware Projection → Semantic Validation Layer → Semantic Controls Layer → FQE → result assembly → lineage record write. Capability availability is declared in the MCP manifest; a capability not enabled by a feature flag or accessible to the user's role appears as `available: false` with a reason.

Example

A structured tool call arrives from AI consumers.

A structured `run_analytics` tool call arrives from the AI Chat Platform. The MCP Capability Layer validates the JWT signature, confirms the token has not expired, and extracts the claims. For a natural language query it routes to the Intent Resolution Agent; for a structured call it routes directly to the Role-Aware Projection Layer, which makes the entitlement decisions before the Semantic Validation Layer compiles the plan. The MCP Capability Layer does not interpret the parameters or make any analytical decisions; it validates, routes, and waits.


Intent Resolution Agent (IRA)

Governing principles: P2 — Controls before execution · P10 — Deterministic computation, not generation

The Intent Resolution Agent is the AI component responsible for translating a natural language query into a structured, validated operation request. It is the only AI step in the pre-computation pipeline. Its output — a resolved `operation_id` and bound `params` — is handed, with the caller's JWT, to the Role-Aware Projection Layer for entitlement decisions before the Semantic Validation Layer compiles the plan.

The IRA does not interpret data, make recommendations, or produce output visible to the end user. Its sole function is operation selection and parameter binding. Once it has resolved intent, it hands off a deterministic structured request and plays no further role in the pipeline.

Intent Resolution Pipeline



RAG Retrieval

At registration time, each `analytical_operation` and `analytical_metric` definition in the SMR is embedded: the operation name, description, example phrasings, required parameters, and associated metric descriptions are concatenated and encoded as a dense vector stored alongside the definition.

When a query arrives, the IRA encodes the natural language input and performs a vector similarity search against the SMR operation embeddings. The top-K candidate operations (and their associated metric definitions) are retrieved. Only these candidates — not the full catalogue — are injected into the LLM ranking prompt.

LLM Intent Ranking

The LLM receives the top-K candidates and the user's query. It ranks candidates, binds parameters, and scores confidence for each. It also derives a `presentation_hint` for each candidate — the likely chart type and primary axes — based on the operation's result shape and the SMR operation definition. This preview is included in every candidate card returned to the consumer.

If the top candidate's confidence score exceeds `intentConfidenceThreshold` (configurable, default 0.75) and leads the second candidate by more than `intentConfidenceBand` (configurable, default 0.1), the IRA proceeds directly to the RAPL with no card shown. Otherwise, up to three ranked candidate cards are returned for the user to select or refine.

The LLM call is constrained: the prompt contains only the candidate operation definitions and the user's query. The LLM has no access to result data, SMR governance metadata, or user entitlements — those are enforced downstream by SVL and RAPL.

Multi-Candidate Selection

When intent is ambiguous, the IRA returns up to three ranked candidate cards in a single `candidates[]` array. Each card represents one plausible interpretation of the user's query, ordered by confidence. The consumer renders all candidates simultaneously, allowing the user to select the one that matches their intent rather than guessing from a single interpretation.

The number of candidates returned is determined by confidence clustering:

Situation	Cards shown
Top candidate above threshold, clear leader	0 — proceeds directly to RAPL
Top candidate below threshold, or top two within confidence band	2 candidates
Top three candidates within confidence band	3 candidates
<code>requiresIntentConfirmation: true</code> on the operation (governance override)	1 candidate — approval required regardless of confidence

The consumer re-submits with `"selected_candidate": <index>` (0-based) to indicate which card the user chose. The IRA then forwards the selected candidate's resolved intent to the RAPL.

Conversational Refinement

After candidate cards are presented, the user may respond with a natural language adjustment rather than selecting a card — "actually make it monthly, not quarterly" or "add tracking error as a metric." The consumer forwards the refinement turn to the IRA alongside the current candidate state. The IRA treats the refinement as a constrained update: it re-runs the LLM with the selected or leading candidate as context and applies the requested changes to that candidate's parameters. An updated card set is returned.

This creates a pre-execution dialogue loop between the user and the IRA:

1. User sends a query → IRA returns candidate cards
2. User refines ("weekly not quarterly") → IRA updates and returns revised cards
3. User selects a card → IRA forwards resolved intent to RAPL → SVL → execution proceeds

The loop is bounded: a maximum number of refinement turns is configurable (`intentRefinementMaxTurns` , default 5). After the limit, the IRA requires the user to select from the current candidate set or start a new query. No data is accessed and no query is executed during the refinement loop — it is entirely within the IRA's pre-execution scope.

Each refinement turn is recorded in the session context and included in the lineage record's `intent_session` field, providing a complete audit trail of how the final resolved intent was reached.

Presentation Preview

Each candidate card includes a `presentation_hint` block derived from the operation's result shape and the SMR operation definition. This gives the user a preview of how the result will be presented before committing to execution — surfacing whether the result will be a grouped bar chart, a line chart, a heatmap, or a table, and which fields will appear on each axis.

The `presentation_hint` is a pre-execution estimate. The DVL produces the authoritative display specification after execution, based on the actual result shape. If the result shape differs from the estimate (for example, because entitlement projection reduces the metric set), the DVL governs. The hint is informational only.

Field	Description
<code>chart_type</code>	Predicted chart type — <code>bar</code> , <code>line</code> , <code>heatmap</code> , <code>scatter</code> , <code>table</code>
<code>primary_dimension</code>	The field that will appear on the X axis or as the primary grouping
<code>measures</code>	Metric fields that will appear as Y axis values or colour encoding
<code>series_by</code>	Dimension used to create series or colour bands (if applicable)

Structured API Path

Consumers that construct explicit `operation_id` + `params` (agentic pipelines, custom analytics UIs, integration tests) can call `run_analytics` with a structured payload directly. MCP routes these calls directly to the RAPL — which makes the entitlement decisions — and then on to the SVL, bypassing the IRA. The `list_operations` tool returns the full entitled operation catalogue for consumers that build their own operation selection UI or inject the catalogue into their own model context.

Example

Running example — portfolio manager asks:

"Show me portfolio returns versus benchmark for my equity portfolios this quarter."

The IRA encodes this query and retrieves the top-3 candidate operations from the SMR:

`compare_portfolios` (score 0.91), `portfolio_summary` (score 0.67), `benchmark_attribution` (score 0.61).

The top candidate exceeds the confidence threshold and leads by more than 0.1. The LLM binds the resolved intent and derives a presentation preview:

```
{
  "operation_id": "compare-portfolio",
  "params": {
    "port
```

Confirm Analytical Query

INTERNAL

Compare Portfolios

Review the resolved intent before execution proceeds.

OPERATION	compare_portfolios
METRICS	portfolio_return · benchmark_return
DIMENSIONS	portfolio_id · asset_class
TIME PERIOD	Quarter to date
FILTER	asset_class = EQUITY
PERFORMANCE IMPACT	620 / 1,000 units
CLASSIFICATION	INTERNAL

Approving will execute the query. Rejecting returns you to the conversation.

Reject

Approve & Execute


```
f  
o  
l  
i  
o  
—  
i  
d  
s  
"  
:  
  
[  
"  
G  
L  
O  
B  
—  
E  
Q  
—  
O  
P  
P  
"  
,  
"  
U  
K  
—  
C  
O  
R  
E  
—  
I  
N  
C  
"  
],  
,  
  
"  
m  
e  
t  
r  
i  
c  
s  
"  
:
```

```
[ "portfolio_return", "benchmark_return" ],  
  "time_period":
```

```
"
q
u
a
r
t
e
r
-
t
o
-
d
a
t
e
"
,

"
f
i
l
t
e
r
s
"
:
[
{
"
d
i
m
e
n
s
i
o
n
"
:
"
a
s
s
e
t
-
c
l
a
s
s
"
,

```

```
"operator":  
{"eq":  
  ,  
  "value":  
    :  
    "EQUITY"  
  }  
}  
 ,  
  "presentation - hint":  
  :
```

```
{  
  
  "chart-type":  
  
  "bar",  
  
  "primary-dimension":  
  
    "portfolio"
```

```
-id",  
    "measures":
```

```
[  
    "portfolio-return",  
    "benchmark-return"
```

```
rn"  
]  
,  
  
"  
s  
e  
r  
i  
e  
s  
-  
b  
y  
":  
  
"  
m  
e  
t  
r  
i  
c  
"  
  
}  
}
```

Confidence is 0.91 — above threshold, no candidate cards shown. Resolved intent is forwarded to the RAPL.

↳ **Step 1 — Intent resolved.** The IRA has translated the natural language query into a structured, bound operation request with a presentation preview. The RAPL now makes the entitlement decisions, then the SVL validates and plans.

Semantic Metrics Repository (SMR)

Governing principles: P1 — Semantic abstraction · P3 — Deterministic metric resolution · P9 — Administrator sovereignty

The Semantic Metrics Repository (SMR) is the governing catalogue of every analytical concept resolvable on the platform. Before any query can be planned or executed, every identifier in that query (metrics, dimensions, hierarchies) must be registered in the SMR. This is an architectural constraint, not a policy: the Semantic Validation Layer rejects any identifier not present in the SMR, and nothing is queryable that is not registered.

Concept Types

The SMR holds three types of JSON metadata definition:

Definition type	Description
<code>analytical_metric</code>	Governed metric definition — formula, aggregation, <code>data_affinity</code> , <code>physical_mapping</code> , <code>required_dimensions</code> , <code>performance_impact_weight</code> , <code>classification_level</code> , <code>compliance_relevant</code> , <code>regulatory_framework</code>
<code>analytical_dimension</code>	Governed dimension definition — <code>data_affinity</code> , <code>physical_mapping</code> , enumerated values or <code>hierarchical_flag</code> , <code>hierarchy_levels</code>
<code>analytical_operation</code>	Governed operation definition — <code>execution_profile</code> , <code>required_params</code> , <code>supported_metrics</code> , <code>supported_dimensions</code> , <code>default_visualization</code>

Metric Definition Schema

Every metric in the SMR conforms to the following schema. This is the authoritative reference for metric registration and validation:


```

{
  "id": "portfolio_return",
  "version": "2.1.0",
  "label": "Portfolio Return",
  "description": "Total return of a portfolio over the specified period, net of fees, expressed as
a percentage.",
  "formula": "(end_market_value - start_market_value + cash_flows) / start_market_value",
  "compliance_relevant": false,
  "regulatory_framework": [],
  "unit": "percentage",
  "aggregation": {
    "default": "value_weighted_average",
    "allowed": ["value_weighted_average", "equal_weighted_average", "sum"],
    "granularity": ["daily", "monthly", "quarterly", "annual", "since_inception"]
  },
  "dimensions": {
    "required": ["portfolio", "date"],
    "optional": ["asset_class", "currency", "benchmark"]
  },
  "data": {
    "domain": "portfolio",
    "sub_domain": "performance",
    "source_tables": ["fact_portfolio_daily", "dim_portfolio"],
    "refresh_cadence": "daily",
    "latency_sla": "T+1"
  },
  "governance": {
    "owner": "head_of_performance_analytics",
    "steward": "performance_analytics_team",
    "classification": "INTERNAL",
    "approved": true,
    "approved_by": "cdo_office",
    "approved_at": "2025-11-15T09:00:00Z",
    "effective_from": "2025-11-15",
    "deprecated": false
  },
  "lineage": {
    "upstream_metrics": [],
    "upstream_sources": ["positions_service", "pricing_service", "cash_flow_service"],
    "downstream_metrics": ["active_return", "information_ratio"]
  },
  "access": {
    "roles": ["portfolio_manager", "risk_officer", "analytics_governance"],
    "public": false
  },
  "display": {
    "format": "percentage",
    "decimals": 2,
    "sign_convention": "positive_is_gain",
    "benchmark_comparison": true
  }
}

```

All metric definitions must pass through a governance review and approval process before they are resolvable on the platform. A metric authored by a Metrics Modeller is not queryable until it has been reviewed and approved by Analytics Governance.

Formula Language

The SMR formula language expresses metric computation logic in terms of other registered metrics or canonical data source identifiers — never physical column names. This decouples metric definitions from backend schema changes: renaming a physical table or column requires only a mapping update, not a metric definition change. The formula language supports arithmetic composition, conditional expressions, safe division with null protection, time-windowed aggregations, and filtered sub-aggregations.

SMR Authoring and Discovery

The SMR and SDR are independent stores, both housed within the Data Context Store (DCS). The SDR contains the organisation's existing data definitions — data models, object models, physical schemas, and data lineage — and will already exist in most organisations before the Analytics Platform is deployed. The SMR is a separate store for metric definitions; it references the SDR's data definitions but does not extend or depend on it structurally. The DCS is the external container that holds both.

Metric authoring uses the DCS's native versioning and approval workflow. There are no custom MCP tools for metric authoring. Metrics Modellers author `analytical_metric`, `analytical_dimension`, and `analytical_operation` JSON metadata definitions through the DCS authoring interface; Analytics Governance approves them before they become resolvable.

Discovery — AI models and agents discover available operations by calling the `list_operations` MCP tool (defined in the MCP Capability Layer). `list_operations` returns operation IDs, display names, required parameters, supported metrics, supported dimensions, and execution profiles for all approved operations within the caller's entitlement scope. There are no `smr://` MCP resource URIs and no separate `list_metrics`, `get_metric_definition`, `propose_metric`, or `approve_metric` MCP tools.

The Analytics Engine queries the SMR directly at request time — the SVL resolves metric and operation definitions, and the FQE resolves physical mappings. The DCS also exposes its own external API and MCP interface through which consumers and tooling can independently browse, inspect, and discover SMR metric definitions and SDR data definitions without going through the Analytics Engine.

Example

Continuing the portfolio manager query — when the request reaches the SVL (Stage 1 validation, downstream of RAPL), the SMR is the catalogue every identifier is resolved against.

The SVL asks the SMR to resolve the `compare_portfolios` operation, then resolves `portfolio_return` and `benchmark_return` as `analytical_metric` documents. The SMR confirms both are approved for the `portfolio_manager` role and returns their definitions, including `data_affinity ( portfolio )`, `required_dimensions ( portfolio_id , time_period )`, and `aggregation ( value_weighted_average )`.

`asset_class` resolves as an approved `analytical_dimension` document with an approved filter operator (`eq`). If either metric document were absent or not in `"status": "approved"` state, the SVL would return `METRIC_NOT_FOUND` and the pipeline would stop here.

Role-Aware Projection Layer (RAPL)

Governing principles: P5 — Role-aware by default · P1 — Semantic abstraction

The Analytics Engine's domain is analytics and data mining — it operates on governed metric definitions, not raw tables. The Role-Aware Projection Layer is the component that determines what each user is entitled to see within that domain. It enforces access at two levels: the **analytical level** (which metrics and dimensions a user may query) and the **data level** (which rows and columns are visible within an entitled result).

This distinction is significant. A user may be restricted from raw row-level or column-specific data and yet be fully entitled to receive governed aggregated results. RAPL manages this boundary at the semantic layer: entitlement is conferred by role membership against registered metric definitions, not by database permissions. A portfolio manager who cannot access the positions table can still receive `portfolio_return` as an aggregated metric because their role grants access to that metric definition, and the row predicates RAPL injects ensure the result covers only their authorised portfolios.

Entitlement policies are managed in the **Data Entitlements Store (DES)** — an independent external component. Policies are defined at the **logical object and data element level**, not at the system or database level. A policy grants or restricts access to a named metric, a named dimension, or a named data element as governed concepts — not to a specific table, schema, or connection string. This means entitlements remain stable as the underlying physical implementation changes, and they are comprehensible to business owners and governance teams who have no visibility into the data platform's internal structure.


Projection is not optional and not bypassable. Every request passes through RAPL. RAPL sits after the Intent Resolution Agent and before the Semantic Validation Layer. It receives the fully qualified analytical request — the resolved operation, metric IDs, dimension IDs, and params — together with the caller's JWT. Only with the complete request can RAPL make YES/NO entitlement decisions per metric and dimension, resolve row predicate templates to the user's actual data scope, and register the column masks that apply. The resulting entitlement projection is passed to the SVL, which enforces it in Stage 3 before any query plan is compiled or any backend is contacted.

Restriction Types

RAPL makes four categories of entitlement decision. Every decision is made inside RAPL (Stage 5 of the projection lifecycle); the resulting conditions are carried in the entitlement projection and enforced downstream — metric and dimension removals by the SVL during plan generation, row and column restrictions by the FQE at execution and result assembly.

Restriction type	Decided in RAPL	Enforced at
Metric access	Stage 5 — requested metric not in the entitled access set is DENIED	SVL Stage 3 — denied metric removed from plan; request rejected if a required metric is lost
Dimension access	Stage 5 — requested dimension not in the entitled access set is DENIED	SVL Stage 3 — denied dimension removed from plan
Row scope predicate	Stage 5–6 — population predicate decided and resolved against JWT claims	FQE physical query generation — injected as a predicate node
Column mask	Stage 5 — masked columns and masking mode registered	FQE result assembly — value replaced, redacted, or excluded

Projection Lifecycle



Stage 1 — JWT Validation. Validate the inbound JWT: signature, expiry, and org claim. **DENY** — request rejected immediately if any check fails. No further processing occurs on an invalid token.

Stage 2 — Role Claim Extraction. Extract the user's analytical role claims from the validated JWT using the configured `roleClaimField` (e.g. `analytics_roles`). **DENY** — if no valid analytical role claims are present the request is rejected. All extracted roles are carried forward for merging.

Stage 3 — DES Role Definition Retrieval. Look up the full role definition for each extracted role from the Data Entitlements Store (DES). Each definition declares the metric access set, dimension access set, row scope predicate templates, column masks, and classification ceiling for that role. **DENY** — if no valid role definitions are retrieved the request is rejected.

Stage 4 — Multi-Role Merge. Merge all retrieved role definitions into a single entitlement profile. Metric and dimension access: union. Row scope predicates: intersection (most restrictive wins). Column masks: union (masked by any role = masked for the user). No APPROVE/DENY decision is made here — this stage produces the profile against which all decisions are made in Stage 5.

Stage 5 — Entitlement Decisions. Four classes of decision are made against the merged entitlement profile:

- **Metrics** — each requested metric is **APPROVED** or **DENIED** (`METRIC_NOT_ENTITLED`). Approval may be with or without dimensional constraints — a role may grant access to a metric only at a summary level (no dimension slicing permitted), only with specific aggregations, or with full dimensional access. If any required metric is denied the request is rejected.
- **Dimensions** — each requested dimension is **APPROVED** or **DENIED** (`DIMENSION_NOT_ENTITLED`). Entitlement is evaluated per metric-dimension combination — a dimension approved for one metric may not be approved for another.

- **Row scope** — each approved metric carries an **APPROVED with population restrictions** decision: the row scope predicates that define which population of data the user is permitted to see within that metric's result — for example, only their managed portfolios, only their assigned entities, only their region.
- **Columns** — each approved metric carries an **APPROVED with column restrictions** decision where applicable: which columns are fully visible, which are masked (value replaced or redacted), and which are excluded entirely from the result.

No approved metric reaches the output without its full set of population scope and column visibility conditions attached.

Stage 6 — Row Scope Template Resolution. Resolve the row scope predicate templates from Stage 5 against the user's JWT claims. `{{user.claim_name}}` syntax is expanded to concrete values at query time, producing the exact data population each approved metric is scoped to for this user. **DENY** — if a required JWT claim for scope resolution is missing the request is rejected.

Stage 7 — Entitlement Projection Output. Produce the completed entitlement projection: approved `metric_access_set` (each with permitted aggregations and dimensional constraints), approved `dimension_access_set`, resolved `row_scope_predicates[]`, registered `column_masks[]`, and `classification_ceiling`. Every entry in the output is either an approval or an approval with restrictions — there are no unexamined metrics, dimensions, populations, or columns in the output. Write the full projection record to the ALS. Pass to SVL for Stage 3 enforcement.

Multi-Role Merge Strategy

Entitlement type	Merge strategy	Rationale
Metric access	Union	Entitlement via any role is sufficient
Dimension access	Union	Entitlement via any role is sufficient
Row scope predicates	Intersection (AND)	All predicates must be satisfied — most restrictive wins
Column masks	Union	A column masked by any role is masked for the user

Column Masking Modes

Three masking modes are supported for column restrictions applied in Stage 5 and enforced by the FQE during result assembly:

Mode	Column representation in result
<code>null_replacement</code>	Column value replaced with <code>null</code>
<code>redacted_label</code>	Column value replaced with <code>"[REDACTED]"</code>
<code>excluded</code>	Column omitted entirely from the result schema

Entitlement Audit

Every entitlement decision — approved metrics, denied metrics, approved dimensions, row scope predicates applied, and column masks registered — is recorded in the Analytical Lineage Store (ALS) as part of the projection record. The record captures the roles active at query time, each metric's APPROVE or DENY outcome and the basis for it, the resolved row scope population for each approved metric, and the column restrictions in force. This record provides the evidentiary chain required for regulatory entitlement audits.

Example

With the operation and metrics resolved, RAPL constructs the entitlement projection.

The RAPL reads the `portfolio_scope` claim from the JWT and resolves it against the `portfolio_manager` role's predicate template (`portfolio_id IN ({user.portfolio_scope})`):

```
{
  "row_predicates": [
    { "dimension": "portfolio_id", "operator": "in",
      "values": ["GLOB_EQ_OPP", "UK_CORE_INC", "ASIA_PAC_GRW", "EUR_BAL_INC"] }
  ],
  "column_masks":      [],
  "classification_ceiling": "INTERNAL",
  "projection_basis":    "jwt_claim:portfolio_scope"
}
```

No column masks apply — the `portfolio_manager` role has no masking rules for performance metrics. The row predicate is injected into the LQP as node `n3` (see SVL example below). Any portfolio outside the four listed IDs is excluded at the physical query level.

↳ **Step 2 — Entitlement projection complete.** Both metrics are confirmed as entitled for the `portfolio_manager` role. The user's access scope is locked — row scope predicates resolved, entitled portfolio set established, column masks registered. The entitlement projection is handed to the SVL. No backend has been contacted.

Semantic Validation Layer (SVL)

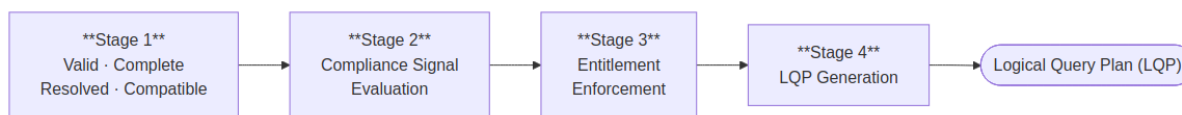
Governing principles: P2 — Controls before execution · P10 — Deterministic computation, not generation

The Semantic Validation Layer receives the entitlement projection from the Role-Aware Projection Layer together with the fully qualified analytical request — either from a structured API caller or as the resolved output of the Intent Resolution Agent — and produces a validated, platform-agnostic Logical Query Plan (LQP). It is entirely deterministic: no AI model runs inside it. Its purpose is to: (1) resolve the operation from the SMR catalogue, (2) validate `params` against the operation's `required_params` schema, (3) resolve metric IDs within `params` against `analytical_metric` metadata definitions, (4) enforce the entitlement projection

computed by RAPL, (5) build the LQP. The output (the LQP) contains no backend references, no SQL, and no physical schema identifiers: only analytical operations expressed against SMR-registered concepts.

Four-Stage Validation Pipeline

Every analytical request passes through four sequential stages:



Stage 1 — Valid, Complete, Resolved & Compatible. The request must be well-formed, fully populated, and resolvable against the SMR before any further processing occurs. Required fields must be present and correctly typed. The `operation_id` must resolve to an approved `analytical_operation` metadata definition. Every metric ID must resolve to an approved `analytical_metric` metadata definition and every dimension ID to an approved `analytical_dimension` metadata definition. Params must conform to the operation's `required_params` schema. Cross-entity compatibility is validated in the same pass: required dimensions must be present for each metric, aggregation rules must be mutually compatible, time granularity must be consistent, and filter predicates must reference valid fields. Anything unregistered, unapproved, missing, or incompatible is rejected here — the pipeline does not proceed.

Stage 2 — Compliance Signal Evaluation. The SVL combines two independent signals to determine the compliance disposition of the request. Signal 1 is the IRA's compliance intent score — derived from the user's natural language query and forwarded as part of the resolved request. Signal 2 is the `compliance_relevant` flag on each resolved `analytical_metric` metadata definition, set by the Metrics Modeller at registration. If both signals are active the request is escalated to the full compliance tier and the Provenance Artifact Service will be invoked. If only the metric signal is active the request proceeds as a standard governed query. If the user's stated intent is compliance-driven but the requested metrics are not registered as compliance-relevant, the SVL rejects the request — the metric definition is not approved for compliance use and the query must not proceed silently with an inappropriate definition.

Stage 3 — Entitlement Enforcement. The SVL applies the entitlement projection computed by the Role-Aware Projection Layer. Metrics and dimensions are filtered to the caller's entitled scope. Row scope predicates resolved by RAPL are injected into the plan. Column masks are registered for enforcement at result assembly. Any metric or dimension RAPL did not approve is removed; if removal leaves the request without its required metrics the request is rejected with an entitlement error rather than returning a partial result.

Stage 4 — LQP Generation. The validated, projected, and compliance-classified request is compiled into a platform-agnostic Logical Query Plan — a directed acyclic graph of analytical operations expressed entirely in SMR-registered concepts. No SQL, no backend references, and no physical schema identifiers appear in the LQP. Data affinity hints are assigned per metric node to guide the Physical Query Planner's sub-plan decomposition. Result cardinality and execution performance impact are estimated and attached to the plan for the Semantic Controls Layer to evaluate before execution is authorised.

Example

The SVL receives the fully qualified analytical request together with RAPL's entitlement projection and passes them through all four stages. The output is the Logical Query Plan passed to the Semantic Controls Layer:

```
{
  "lqp_id": "lqp-20260518-093243-r9xq",
  "intent_id": "int-20260518-093241-pk7m",
  "org_id": "acme-wealth",
  "nodes": [
    { "id": "n1", "type": "metric_scan",
      "metrics": ["portfolio_return", "benchmark_return"],
      "dimensions": ["portfolio_id"],
      "time_period": { "granularity": "quarterly", "anchor": "current" } },
    { "id": "n2", "type": "filter", "input": "n1",
      "predicate": { "field": "asset_class", "operator": "eq", "value": "EQUITY" } },
    { "id": "n3", "type": "filter", "input": "n2",
      "predicate": { "field": "portfolio_id", "operator": "in",
        "values": ["GLOB_EQ_OPP", "UK_CORE_INC", "ASIA_PAC_GRW", "EUR_BAL_INC"] } },
    { "id": "n4", "type": "sort", "input": "n3",
      "field": "portfolio_return", "direction": "desc" }
  ],
  "output_node": "n4",
  "estimated_performance_impact": 620,
  "classification_required": "INTERNAL"
}
```

Node `n3` is the RAPL row scope predicate — part of the plan, not a post-execution filter.

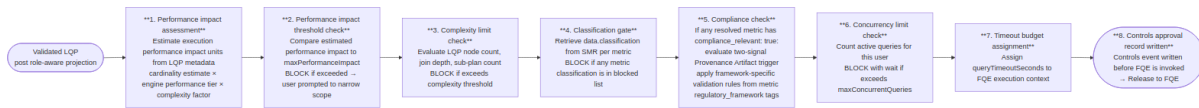
↳ **Step 3 — Validation and planning complete.** The fully qualified request has been validated against the SMR and compiled into a platform-agnostic Logical Query Plan, with RAPL's entitlement projection enforced into the plan. Every metric and dimension identifier is confirmed as registered and approved. No backend has been contacted.

Semantic Controls Layer (SCL)

Governing principles: P2 — Controls before execution · P8 — Explainability at every layer · P9 — Administrator sovereignty

The Semantic Controls Layer (SCL) applies a suite of performance impact thresholds, complexity limits, and compliance classification checks to every query before it is released to the Federated Query Engine. It is the final gate before physical execution. Controls apply to every query without exception. There is no privileged user, trusted agent, or internal path that bypasses SCL checks.

Controls Pipeline



Performance Impact Assessment Model

Performance impact units are estimated from LQP metadata before any backend is contacted:

Factor	Contribution
Number of metrics	<code>metric_count × 50</code> base units
Engine performance tier per sub-plan	<code>minimal: 10, low: 50, standard: 100, high: 300, unrestricted: 0</code>
Dimension cardinality	<code>low: ×1.0, medium: ×1.5, high: ×3.0, unbounded: ×5.0</code>
Time period scope	<code>single_day: ×1.0, quarter: ×2.0, year: ×4.0, since_inception: ×8.0</code>
Number of sub-plans (federation)	<code>+100 per additional sub-plan</code>
Materialised view match	<code>-800 (pre-computed result)</code>
Cache hit (estimated)	<code>-900 (full cache hit expected)</code>

Worked example — `portfolio_return, tracking_error BY portfolio, asset_class FOR YEAR_TO_DATE :`

```

portfolio_return:      50 (metric base)
tracking_error:        50 (metric base)
SQL warehouse backend: 100 (standard performance tier)
semantic layer backend: 50 (low performance tier)
asset_class:           1.5× cardinality multiplier (medium)
YEAR_TO_DATE:          4.0× period multiplier
2 sub-plans:           100 (federation overhead)

```

Base: $(50+50) \times 1.5 \times 4.0 = 600$

Engines: $100+50 = 150$

Federation: 100

Total estimate: 850 performance impact units

Against a `maxPerformanceImpact: 1000` limit, this query is approved. Against a `500` limit, it is blocked and the user receives structured suggestions to narrow scope (reduce time period, reduce metric count, or add a filter).

Compliance

The platform's compliance behaviour is driven entirely by the metrics being queried and the intent of the query. There is no tenant-level compliance mode switch. The regulatory framework is declared on each metric at registration time by the Metrics Modeller.

Feature flag

Compliance features are enabled or disabled at the tenant level with a single binary flag, set by the Platform Admin:

```
"features": {  
  "complianceMode": true  
}
```

When `complianceMode` is `false`, all compliance checks and Provenance Artifact generation are disabled. When `true`, the two-signal trigger below applies to every query.

Provenance Artifact trigger — two signals, both required (AND logic)

The platform escalates to the enhanced Provenance Artifact only when both of the following signals are true at runtime:

Signal	Source	True when
Signal 1 — metric metadata	<code>compliance_relevant</code> field on <code>analytical_metric</code> SMR definition	At least one resolved metric has <code>compliance_relevant: true</code> . Set by the Metrics Modeller at registration.
Signal 2 — AI intent classification	Semantic Validation Layer compliance intent classification (Stage 2b)	<code>compliance_purpose_score</code> $\geq$ the platform-level <code>compliance_intent_threshold</code> (default 0.8, configurable). The SVL classifies the query's analytical intent — derived from the operation ID, resolved metric descriptions, and parameter values — and sets <code>compliance_purpose: true</code> if the score meets the threshold.

Combined decision:

<code>compliance_relevant</code> (any metric)	<code>compliance_purpose</code> (SVL classification)	SCL decision
<code>true</code>	<code>true</code>	Enhanced — full Provenance Artifact active; framework-specific validation rules applied
<code>true</code>	<code>false</code>	Standard controls output
<code>false</code>	<code>true</code>	Standard controls output

compliance_relevant (any metric)	compliance_purpose (SVL classification)	SCL decision
false	false	Standard controls output

Export gate

When the Provenance Artifact is active, export of the result is blocked until the artifact is confirmed written and sealed to the ALS. The `export_requires_lineage: true` flag in the response signals this state to the consumer. The consumer must not present export affordances until the platform confirms sealing.

Framework-specific validation rules

Each `regulatory_framework` tag on a metric carries validation rules that the SCL applies when the Provenance Artifact is active for that metric. These rules are declared at metric registration and enforced at query time — they are properties of the metric, not tenant configuration.

Framework	Additional validation rule	Effect
<code>mifid2</code>	Queries on client-identifiable metrics require business justification	User is prompted for justification before execution; justification recorded in Provenance Artifact
<code>mifid2</code>	Best-execution metrics require explicit timeframe	Validation error if <code>date</code> dimension not specified
<code>mifid2</code>	Transaction reporting metrics generate a regulatory trace record	Additional trace written to the ALS regulatory trace partition
<code>basel3</code>	Capital ratio metrics require entity identifier	Validation error if <code>entity</code> dimension not specified
<code>basel3</code>	LCR/NSFR metrics generate a daily snapshot record	Snapshot written to the ALS regulatory snapshots partition
<code>basel3</code>	Stress scenario metrics are subject to RESTRICTED classification ceiling	Classification ceiling applied regardless of user role
<code>sec_reg_bi</code>	Client advisory metrics: narrative synthesis constrained to factual summary only	Forward-looking statements, yield projections, and investment recommendations are rejected from NSA output before inclusion in the response
<code>sec_reg_bi</code>	Advisory queries require suitability record reference	<code>suitability_record_id</code> parameter required before execution

Trace target routing to the correct ALS partition is determined automatically from the `regulatory_framework` tags on the resolved metrics.

Timeout and Partial Result Handling

Scenario	Behaviour
All sub-plans complete within timeout	Normal result assembly and return
One sub-plan times out, others complete	Partial result assembly — missing metrics represented as null with <code>timeout</code> provenance marker; user notified
All sub-plans time out	Query failed — error returned to user; controls event written with <code>timeout</code> status
Engine cancellation on timeout	FQE sends cancellation signal to timed-out engine (if engine supports cancellation)

Example

The projected LQP reaches the SCL for controls validation.

SCL evaluates the LQP (`lqp-20260518-093243-r9xq`) against the `acme-wealth` controls config:

Check	Value	Limit	Result
Estimated performance impact	620	1,000	Pass
Metrics per query	2	10	Pass
Dimensions	1	5	Pass
Data classification	INTERNAL	INTERNAL ceiling	Pass
Compliance	none required	—	Pass

All checks pass. SCL writes a controls decision record to the Analytical Lineage Store (ALS) — before the FQE is invoked — and releases the LQP to the FQE:

```
{
  "lqp_id": "lqp-20260518-093243-r9xq",
  "decision": "approved",
  "timestamp": "2026-05-18T09:32:44Z",
  "checks": ["performance_impact_ceiling", "metric_count", "dimension_count",
    "classification_gate", "compliance_check"],
  "result": "all_passed"
}
```

Physical Query Planner (PQP)


Governing principles: P1 — Semantic abstraction · P10 — Deterministic computation, not generation

The Physical Query Planner (PQP) is the translation boundary between the logical and physical layers of the pipeline. It receives the controls-approved Logical Query Plan from the SCL and produces backend-specific physical query fragments — one per data affinity — that the Federated Query Engine can route directly to execution backends.

Nothing above the PQP has knowledge of physical schemas, table names, or backend query languages. Nothing below it operates on logical concepts. The PQP is the single point where semantic intent becomes executable instruction.

Responsibilities

The PQP performs three operations in sequence:



The PQP has no execution capability. It does not connect to backends, manage timeouts, or assemble results. Its output is a set of ready-to-execute physical query fragments; the FQE owns everything from that point forward.

Translation is deterministic

The same LQP always produces the same physical sub-plans. Given identical metric definitions, filter predicates, and time expressions, the PQP's output is fully reproducible. This property is required for lineage: the physical sub-plans written to the lineage record must faithfully represent what was executed, with no runtime variation.

physicalMapping resolution

The `physical_mapping` field on each `analytical_metric` SMR definition declares where the metric lives physically:

Metric field	Resolved to
<code>physical_mapping.source</code>	Registered backend ID — matched against the FQE backend registry
<code>physical_mapping.table</code>	Physical table or view name in the target backend
<code>physical_mapping.measure</code>	Column or pre-computed measure identifier

Metric field	Resolved to
<code>physical_mapping.cube</code>	Cube name for semantic layer backends

These fields are already attached to the metric nodes in the LQP by the SVL at resolution time. The PQP reads them directly — no additional SMR call is required.

Example

The SCL has approved the LQP for the portfolio manager query. The PQP receives it and begins translation.

Both `portfolio_return` and `benchmark_return` carry `data_affinity: "portfolio"` and `physical_mapping.source: "primary-warehouse"`, so the PQP produces a single sub-plan. It reads the physical table and measure references from the metric nodes, applies the RAPL row predicate and asset class filter, expands `quarter_to_date` to the concrete date range `2026-04-01 → 2026-06-30`, and translates the result into SQL:

```
-- Physical sub-plan: affinity=portfolio · backend=primary-warehouse
SELECT
  p.portfolio_id,
  SUM(f.portfolio_return * f.market_value) / SUM(f.market_value) AS portfolio_return,
  SUM(f.benchmark_return * f.market_value) / SUM(f.market_value) AS benchmark_return
FROM fact_portfolio_daily f
JOIN dim_portfolio p ON f.portfolio_id = p.portfolio_id
WHERE p.asset_class = 'EQUITY'
  AND f.portfolio_id IN ('GLOB_EQ_OPP', 'UK_CORE_INC', 'ASIA_PAC_GRW', 'EUR_BAL_INC')
  AND f.date BETWEEN '2026-04-01' AND '2026-06-30'
GROUP BY p.portfolio_id
ORDER BY portfolio_return DESC
```

If the query also included a risk metric with `data_affinity: "risk_metrics"`, the PQP would produce a second sub-plan — translated to the semantic layer's MetricFlow query format — and hand both to the FQE for parallel execution.

↳ **Step 4a — Physical query planning complete.** The approved LQP has been resolved against physical mappings and translated into one executable SQL sub-plan. The FQE will route it to the registered backend. No backend has been contacted yet.

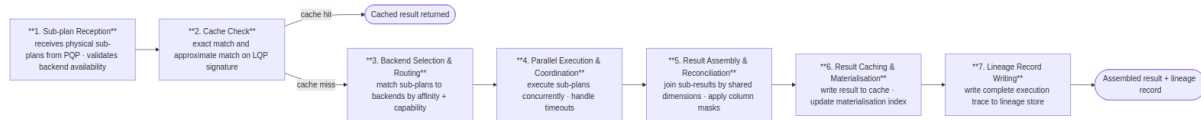
Federated Query Engine (FQE)

Governing principles: P1 — Semantic abstraction · P4 — Complete analytical lineage · P10 — Deterministic computation, not generation

The Federated Query Engine (FQE) is the only component in the platform with knowledge of execution backend connection details — endpoints, credentials, and availability. It receives physical sub-plans from the Physical Query Planner, routes them to the registered execution backends in parallel, assembles the results,

and writes a complete execution record to the lineage store. Sub-plan decomposition and dialect translation are the PQP's responsibility; the FQE owns everything from routing onward.

Nine-Step FQE Pipeline



Backend Routing

The FQE receives physical sub-plans from the PQP — already decomposed by data affinity and translated to the backend's native dialect. It matches each sub-plan to a registered execution backend by affinity and capability, then executes all sub-plans concurrently. Shared dimensions across sub-plans become the join keys for result assembly.

For a query producing two sub-plans — `portfolio_return` (affinity: `portfolio`, SQL) and `var_95` (affinity: `risk_metrics`, MetricFlow) — the FQE routes each to its registered backend, executes them concurrently, and joins the results in memory on `portfolio_id` and `date`.

Backend Adapter Table

Backend type	Protocol	Typical use
SQL warehouse	database connector protocol, SQL dialect	Primary performance and position data
Semantic layer	REST/JSON query API	Pre-modelled metrics via semantic layer backend
OpenData API	REST data API	Reference data and third-party feeds
Graph Data API	graph query API	Relationship and counterparty data
OLAP engine	REST/JSON cube query	Pre-aggregated dimensional data
Custom adapter	Platform adapter SDK	Proprietary or specialised data sources

When multiple engines are registered for the same data affinity, the FQE selects the highest-priority available engine. If the highest-priority engine is unavailable or its p95 latency exceeds twice its baseline over a rolling one-hour window, the FQE automatically routes to the next registered engine for that affinity.

Caching

The FQE maintains a result cache keyed by the LQP signature — a deterministic SHA-256 hash of the metric IDs and versions, dimension IDs, filter predicates, time expression, entitlement hash, and org ID.

Cache property	Specification
Cache key	SHA-256 of (metric IDs + versions, dimension IDs, filter predicates, time expression, entitlement hash, org ID). Entitlement hash is a SHA-256 of the fully resolved <code>row_predicates</code> and <code>column_masks</code> from the RAPL projection record for the request, computed after role merging. Two users with different effective predicates always produce different entitlement hashes and are never served each other's cached results.
Cache TTL	Configurable per <code>data.refresh_cadence</code> in the metric definition. Default: 3600 seconds.
Cache invalidation	On metric definition version change; on execution backend data refresh signal; on explicit cache clear via Admin API
Cache scope	Platform-scoped. All results are isolated to the single deployed organisation.
Cache storage	Platform-managed result cache. Results over 10 MB bypass the cache and are streamed directly.
Cache hit disclosure	Cache hits are disclosed in the lineage record and optionally surfaced to the user as a "Result from cache (data as of [timestamp])" indicator.
Cache bypass	Queries with <code>compliance_purpose: true</code> bypass the cache. Compliance artifacts must be freshly generated for each compliance-purpose execution; cached results cannot retroactively produce regulatory trace records.

Adaptive Planning

The FQE adapts routing decisions based on observed execution performance. It tracks p50/p95 latency per engine per data affinity over a rolling one-hour window, automatically falls back to the next available engine if performance degrades, and calibrates performance impact estimates based on observed execution data from completed queries. If a sub-plan engine returns a partial result due to timeout, the FQE logs this in the lineage record and surfaces a warning to the user alongside the partial result.

Example

The approved LQP is released to the FQE for physical execution.

Both `portfolio_return` and `benchmark_return` have `data_affinity: "portfolio"`, so the FQE routes a single sub-plan to the primary SQL warehouse:


```
-- Physical execution (FQE output)
SELECT
    p.portfolio_id,
    SUM(f.portfolio_return * f.market_value) / SUM(f.market_value) AS portfolio_return,
    SUM(f.benchmark_return * f.market_value) / SUM(f.market_value) AS benchmark_return
FROM fact_portfolio_daily f
JOIN dim_portfolio p ON f.portfolio_id = p.portfolio_id
WHERE p.asset_class = 'EQUITY'
    AND f.portfolio_id IN ('GLOB_EQ_OPP', 'UK_CORE_INC', 'ASIA_PAC_GRW', 'EUR_BAL_INC')
    AND f.date BETWEEN '2026-04-01' AND '2026-06-30'
GROUP BY p.portfolio_id
ORDER BY portfolio_return DESC
```

Assembled result:

```
{
  "result_id": "res-20260518-093247-wk4n",
  "latency_ms": 1187,
  "backends_used": ["primary-warehouse"],
  "rows": [
    { "portfolio_id": "GLOB_EQ_OPP", "portfolio_return": 4.21, "benchmark_return": 3.85 },
    { "portfolio_id": "ASIA_PAC_GRW", "portfolio_return": 3.67, "benchmark_return": 3.90 },
    { "portfolio_id": "UK_CORE_INC", "portfolio_return": 2.87, "benchmark_return": 2.54 },
    { "portfolio_id": "EUR_BAL_INC", "portfolio_return": 1.93, "benchmark_return": 2.31 }
  ]
}
```

The FQE writes an execution record to the Analytical Lineage Store (ALS) and passes the assembled result in parallel to the Data Visualization Language (DVL) and Narrative Synthesis Agent.

↳ **Step 4b — Execution complete.** The physical sub-plans were routed to the registered warehouse, executed, and the result assembled. The full audit record is written. No physical schema was exposed at any stage.

Data Visualization Language (DVL)

Governing principles: P7 — Deterministic visualisation

The Data Visualization Language (DVL) is the governing schema that maps result characteristics and analytical intent patterns to specific, parameterised chart contracts. It exists to make chart selection deterministic: the same analytical pattern produces the same chart type across all users, sessions, and AI model versions, regardless of how the question was phrased. The AI model does not select chart types. Intent signals from the query are treated as inputs to the ontology evaluation algorithm, but DVL makes the final binding decision.

Intent Pattern Taxonomy

Every analytical result is classified into one of seven intent patterns:

Intent pattern	Description	Typical trigger phrases
COMPARISON	Comparing a metric across discrete categories	"compare", "by", "across", "ranked by", "top N"
TREND	Showing how a metric changes over time	"over time", "trend", "history", "30-day", "since"
DISTRIBUTION	Showing the spread or concentration of a metric	"distribution", "breakdown", "composition", "concentration"
THRESHOLD	Comparing a metric against a limit or benchmark	"exceeding", "within limit", "versus benchmark", "breach"
ATTRIBUTION	Decomposing a metric into contributing factors	"attribution", "contribution", "drivers", "breakdown by"
RELATIONSHIP	Showing correlation or dependency between metrics	"versus", "correlation", "scatter", "risk-return"
COMPOSITION	Showing part-to-whole relationships	"proportion", "weight", "allocation", "share"

Chart Contract Table

Contract name	Intent patterns matched	Chart type	Key axis assignments	Interaction semantics
BAR_MULTI_SERIES_COMPARISON	COMPARISON	Bar	X: primary categorical dimension (sorted by primary metric DESC); Y: metric value; Colour: metric series or secondary dimension	Click: drilldown; Hover: tooltip; Selection: multi-point
LINE_TIME_SERIES_TREND	TREND	Line	X: temporal dimension; Y: metric value; Colour: metric series; Reference line: injected if <code>compare_to</code> present	Hover: crosshair tooltip; Click data point: surface lineage; Brush: zoom X-axis
HEATMAP_THRESHOLD_MATRIX	THRESHOLD	Heatmap	X: first categorical dimension; Y: second categorical dimension; Colour: metric as % of threshold (diverging scale, midpoint at 100% of limit)	Click cell: drilldown into dimension intersection

Contract name	Intent patterns matched	Chart type	Key axis assignments	Interaction semantics
TREEMAP_COMPOSITION	COMPOSITION , DISTRIBUTION	Treemap	Area: proportional to metric value; Colour: secondary metric (diverging scale); Label: dimension value + metric value	Click tile: drilldown to next hierarchy level; Hover: tooltip with all metrics
WATERFALL_ATTRIBUTION	ATTRIBUTION	Waterfall	X: contribution dimension; Y: contribution value (positive/negative); Colour: positive (green), negative (red), total (grey)	Hover: contribution value and percentage
SCATTER_RISK_RETURN	RELATIONSHIP	Scatter	X: first metric (conventionally risk); Y: second metric (conventionally return); Colour: categorical dimension; Size: optional third metric	Hover: all metric values; Reference lines: quadrant boundaries from benchmark values if present
TABLE_GOVERNED	Fallback (any)	Table	All result columns; column labels from SMR; inline sparklines for temporal metrics	Column sorting; column filtering; export to CSV and JSON

Priority-Ordered Evaluation Algorithm

The ontology evaluator receives the LQP, intent pattern classification, and result schema. It evaluates contracts in order of specificity and returns the highest-scoring match:

```
def evaluate_ontology(lqp, intent_pattern, result_schema, allowed_charts):
    candidates = []
    for contract in ONTOLOGY_CONTRACTS:
        if not all(c in allowed_charts for c in [contract.chart_type]):
            continue
        score = contract.match_score(lqp, intent_pattern, result_schema)
        if score > 0:
            candidates.append((score, contract))
    candidates.sort(key=lambda x: x[0], reverse=True)
    if candidates:
        return candidates[0][1]
    else:
        return TABLE_GOVERNED # deterministic fallback
```

The `TABLE_GOVERNED` contract is the unconditional fallback. It is always eligible and always matches — ensuring that every query, regardless of result shape, receives a valid `display_spec`.

Override Mechanism

Power Analysts may override the ontology's chart selection for a single result by expressing an explicit chart type preference in their query. Overrides are subject to the requested chart type being in the platform's configured `allowedChartTypes` list and the result schema being compatible with the requested chart type. Incompatible overrides are rejected with an explanation. All overrides are logged in the lineage record as analyst-requested deviations from the governing ontology.

Example

With the assembled result available, the Data Visualization Language (DVL) selects the presentation contract.

The ontology evaluator classifies the result: two metrics across four named entities, with a natural comparison relationship between return and benchmark. This matches the `COMPARISON` pattern. The highest-scoring contract is `BAR_MULTI_SERIES_COMPARISON` — a grouped bar chart with portfolios on the X axis and return values on Y, coloured by metric series.

The ontology produces the following DVL display specification:

```
{
  "type": "chart",
  "mark": "bar",
  "data": { "name": "result" },
  "transform": [
    { "fold": ["portfolio_return", "benchmark_return"], "as": ["metric", "value"] }
  ],
  "encoding": {
    "x": { "field": "portfolio_id", "type": "nominal", "title": "Portfolio" },
    "y": { "field": "value", "type": "quantitative", "title": "Return (%)", "axis":
{ "format": ".2f" } },
    "color": { "field": "metric", "type": "nominal", "title": "Metric",
      "scale": { "domain": ["portfolio_return", "benchmark_return"],
        "range": ["#0057B8", "#A8C8F0"] } },
    "xOffset": { "field": "metric", "type": "nominal" }
  },
  "title": "Portfolio Return vs Benchmark – Q2 2026"
}
```

Narrative Synthesis Agent (NSA)

Governing principles: P6 — Governed narrative · P10 — Deterministic computation, not generation

The Narrative Synthesis Agent (NSA) is a secondary AI component inside the Analytics Engine. It runs after the computation pipeline completes — after the FQE has assembled the result, in parallel with the Data

Visualization Language (DVL), — making a single, tightly-scoped call to a language model with one purpose: summarise the structured result in plain language, anchored strictly to the computed values.

The NSA is not a general-purpose AI model with access to the query, the user's intent, or the SMR. Its prompt is constructed entirely from the assembled result: metric labels, row values, units, and dimension names. It is told what the data shows; it is not told what the user asked. This constraint is intentional — it prevents the narrative from interpreting, recommending, or inferring beyond what was computed.

Anchoring and Validation

The NSA produces two output fields:

Field	Description
<code>narrative.lead</code>	A single sentence summarising the top-level finding (e.g. "3 of your 4 equity portfolios outperformed their benchmark this quarter.")
<code>narrative.detail</code>	Two to four sentences covering the most significant data points, one per dimension value or notable comparison
<code>narrative.anchoredTo</code>	An array of dimension value identifiers from the result that the narrative references — used for post-generation validation

After generation, the NSA runs a validation pass: every numeric value in the narrative must be present in the result set. If any value in the narrative cannot be matched to a result row, the narrative is rejected and a single regeneration is attempted. If the second attempt also fails validation, the `narrative` field is omitted from the response and the failure is recorded in the lineage record. This constraint enforces P6 (governed narrative) at the component level.

Model Selection

Query type	Model	Rationale
Standard queries (≤ 5 metrics, ≤ 3 dimensions)	Standard language model	Low latency; sufficient for concise, anchored summaries
Complex queries (attribution, multi-portfolio, regulatory)	Extended language model	Richer result structure requires more precise prose

The model used is recorded in `narrative_status` in the lineage record.

Feature Flag

Narrative synthesis is controlled by the `features.narrativeSynthesis` platform configuration flag. When disabled, the NSA is not invoked and no `narrative` field is included in the response. The default is enabled. Disabling narrative synthesis has no effect on computation, lineage, or display spec generation.

Example

In parallel with the Data Visualization Language (DVL), the NSA receives the assembled result.

The portfolio manager's query returns four rows. The NSA receives the assembled result and produces:

```
{
  "narrative": {
    "lead": "2 of your 4 equity portfolios outperformed their benchmark this quarter.",
    "detail": "Global Equity Opportunities returned 4.21% against a benchmark of 3.85%. UK Core Income returned 2.87% against its benchmark of 2.54%. Asia Pacific Growth and EUR Balanced Income underperformed, returning 3.67% and 1.93% respectively against benchmarks of 3.90% and 2.31%.",
    "anchoredTo": ["GLOB_EQ_OPP", "UK_CORE_INC", "ASIA_PAC_GRW", "EUR_BAL_INC"]
  }
}
```

Post-generation validation confirms every verbatim numeric value cited in the narrative (4.21, 3.85, 2.87, 2.54, 3.67, 3.90, 1.93, 2.31) is present in the assembled result rows. Validation matches on exact numeric literals extracted from the narrative text — rounding differences or proportional expressions (e.g. "roughly 4.2%") may not be caught; residual hallucination risk applies to non-literal claims. Validation passes. The narrative is included in the MCP response alongside `display_spec` and `data`.

↳ **Step 5 — Presentation decision complete.** The DVL has selected the governed display contract for this result shape and intent pattern. The NSA has produced a plain-language summary anchored strictly to the computed values. The response is ready to return to the AI consumer.

Analytical Lineage Store (ALS)

Governing principles: P4 — Complete analytical lineage · P8 — Explainability at every layer

The Analytical Lineage Store provides computation provenance: a complete, queryable record of how every result was calculated. Analytical lineage, as defined on this platform, is distinct from data lineage. Data lineage tracks how data moves between systems. Analytical lineage records how the analytics engine used specific metric definitions, entitlement rules, and execution backends to compute a specific result. The lineage record is not a log — it is a first-class data structure. A regulator, auditor, or internal reviewer must be able to reconstruct exactly how a specific number was calculated, by whom, under what entitlements, from which backends, and with what result — without re-running the query.

Storage Design

Lineage records are stored in an object store — one JSON document per query at key `lineage/{org_id}/{yyyy}/{mm}/{dd}/{result_id}.json`. Records are write-once and never mutated. Post-hoc compliance annotations are written as sibling documents (`{result_id}_amendment_{n}.json`) referencing the original `result_id`.

A thin relational database search index holds only scalar fields required for filtered search queries. The full record is always fetched from the object store; the index is never the source of truth for record content.

Per-Query Stored Elements

Element	Storage	Content
Lineage record	Object store — <code>lineage/{org_id}/{yyyy}/{mm}/{dd}/{result_id}.json</code>	Complete chain: tool call parameters → SMR resolution → projection record → LQP → controls decision → FQE execution record → result schema → visualisation contract → narrative synthesis status
SMR snapshot	Embedded in lineage record (<code>resolved_metrics</code>)	For each metric in the query: metric ID, SMR definition version at query time
Projection record	Embedded in lineage record	Roles, requested metrics, projected metrics, blocked metrics, row predicates, column masks
FQE execution record	Embedded in lineage record (<code>sub_plans</code>)	Sub-plan details, engine IDs, latencies, performance impact units, cache hit status
Controls decision	Embedded in lineage record (<code>controls_decision</code>)	Threshold decisions, classification gates, performance impact limit checks — including blocked queries
Search index row	Relational database <code>analytics.lineage_index</code>	Scalar fields for filtered search — <code>result_id</code> , <code>org_id</code> , <code>user_sub</code> , <code>regulatory_frameworks</code> , <code>error_code</code> , <code>cache_hit</code> , <code>created_at</code> , <code>expires_at</code>
Result artefact	Object storage	CSV result set, chart SVG, narrative text — stored per query

Search Index DDL: `analytics.lineage_index`

```
CREATE TABLE analytics.lineage_index (
  result_id      TEXT          PRIMARY KEY,
  org_id         TEXT          NOT NULL,
  user_sub       TEXT          NOT NULL,
  regulatory_frameworks TEXT,
  error_code     TEXT,
  cache_hit      BOOLEAN      NOT NULL,
  created_at     TIMESTAMPTZ  NOT NULL,
  expires_at     TIMESTAMPTZ  NOT NULL
);
```

The index is used only for search and filter queries (finding all records for a user, within a time window, by regulatory framework, etc.). Filtered search results are resolved to full records by fetching each document from the object store using its `result_id` .

Data Isolation

Every lineage document is stored under an `org_id`-prefixed key in the object store (`lineage/{org_id}/...`), and every row in the `analytics.lineage_index` table carries an `org_id` column. Row-Level Security (RLS) in the relational database enforces access isolation on the index table. Object store access is gated by the platform's API, which validates the JWT `org_id` claim before resolving any object key.

Retention

Rule	Specification
Query records	Platform default: 2,555 days (7 years) — covering most regulatory audit look-back periods. Configurable.
Lineage records	Retained at least as long as the corresponding query record. Cannot be deleted independently.
SMR metric versions	Retained indefinitely — metric version history must be preserved for lineage reconstruction.
Controls events	Retained at least as long as query records.
Result artefacts (object storage)	Default: 365 days. Configurable. Lineage record references are preserved even after object storage expiry.
Blocked queries	Retained in full — queries that fail governance checks are as important to retain as successful ones.

Immutability

Lineage records are never modified after writing. There is no update or delete path available to any user — including Platform Admin. Corrections to erroneous records produce new records that reference the original via a `supersedes` relationship. This constraint is enforced at the database layer, not only by application logic.

Regulatory Audit Export

```
GET /v1/audit/queries
  ?user_id={user_id}
  &from={ISO8601}
  &to={ISO8601}
  &metric_ids[]={metric_id}
  &include_lineage=true
→ Returns all queries matching the filter with full lineage records
```

Audit export packages include query records with timestamps and user identifiers, lineage records with metric definition versions, controls decisions, role-aware projection records showing entitlements in force at query time, and result artefacts within the retention period. All export packages are digitally signed by the platform using the platform signing key registered at deployment.

Example

Two lineage writes occur for this query — one before execution, one after.

The first is written by SCL before the FQE is invoked — capturing the controls decision regardless of whether execution succeeds. The second is written by the FQE after execution, producing the full lineage document at

`lineage/acme-wealth/2026/05/18/res-20260518-093247-wk4n.json` .

Controls decision record (written to object store before FQE execution):

```
{
  "result_id": "res-20260518-093247-wk4n",
  "org_id": "acme-wealth",
  "lqp_id": "lqp-20260518-093243-r9xq",
  "event": "controls_approved",
  "timestamp": "2026-05-18T09:32:44Z",
  "checks": ["performance_impact_ceiling", "metric_count", "dimension_count",
    "classification_gate", "compliance_check"],
  "result": "all_passed"
}
```

Full lineage document (written to object store after FQE completes, at `lineage/acme-wealth/2026/05/18/res-20260518-093247-wk4n.json`):

```
{
  "result_id": "res-20260518-093247-wk4n",
  "org_id": "acme-wealth",
  "user_sub": "idp|user_xyz",
  "lqp_id": "lqp-20260518-093243-r9xq",
  "cache_hit": false,
  "controls_decision": {
    "approved": true,
    "checks_passed": ["performance_impact_ceiling", "metric_count", "dimension_count",
"classification_gate", "compliance_check"]
  },
  "sub_plans": [
    {
      "backend": "primary-warehouse",
      "dialect": "sql",
      "latency_ms": 1187,
      "row_count": 4,
      "cache_hit": false
    }
  ],
  "resolved_metrics": [
    { "metric_id": "portfolio_return", "version": "2.1.0" },
    { "metric_id": "benchmark_return", "version": "1.4.2" }
  ],
  "projection_record": {
    "roles": ["portfolio_manager"],
    "row_predicates": ["portfolio_id IN
('GLOB_EQ_OPP', 'UK_CORE_INC', 'ASIA_PAC_GRW', 'EUR_BAL_INC')"],
    "column_masks": []
  },
  "visualisation": { "contract": "BAR_MULTI_SERIES_COMPARISON" },
  "narrative_status": { "generated": true, "validation": "passed" },
  "created_at": "2026-05-18T09:32:47Z",
  "expires_at": "2033-05-18T09:32:47Z"
}
```

The document is immutable from the moment of writing. A corresponding row is inserted into `analytics.lineage_index` for search access. The full chain — original query, controls decision, metric definition versions, entitlements, physical backend call, and chart contract — is retrievable from the object store under `result_id: res-20260518-093247-wk4n`.

Provenance Artifact Service (PAS)

Governing principles: P4 — Complete analytical lineage · P2 — Controls before execution · P9 — Administrator sovereignty

The Provenance Artifact Service (PAS) assembles the Provenance Artifact and returns the sealed compliance block to the MCP tool response. It is invoked in parallel with DVL and NSA, but only when the Semantic Controls Layer has determined that the two-signal compliance trigger is active (`compliance_tier.active: true`). For all other queries the PAS is not invoked and no compliance block appears in the response.

Purpose

The Provenance Artifact is a sealed, tamper-evident record that satisfies regulatory requirements for demonstrable audit trail on compliance-purpose queries. It is distinct from the standard lineage record that every query produces. The lineage record is the full computation trace. The Provenance Artifact is the governed, signed output of that trace — structured for regulatory consumption, explicitly linked to the frameworks that triggered it, and sealed before the result is returned to the consumer.

Assembly and sealing

The PAS reads the controls decision record and execution record that the ALS has written for the current query. It assembles them into a Provenance Artifact document — adding regulatory trace identifiers, framework tags, and the compliance intent score — and seals it by writing the sealed artifact back to the ALS as a sibling document (`{result_id}_provenance.json`). The artifact is immutable from the moment of sealing. No further write or amendment is permitted without creating a new amendment document referencing the original.

Export gate

Until sealing is confirmed, export of the query result is blocked. The PAS sets `export_requires_lineage: true` in the compliance block it returns to the MCP layer. The MCP layer includes this flag in the tool response; the consumer must not present export affordances until the platform confirms sealing is complete.

Compliance block structure

The PAS returns a structured compliance block to the MCP layer for inclusion in the tool response:

Field	Description
<code>compliance_purpose</code>	<code>true</code> — confirms the two-signal trigger was active for this query
<code>intent_score</code>	The <code>compliance_purpose_score</code> from SVL — the signal that crossed the <code>compliance_intent_threshold</code>
<code>triggered_by_metrics</code>	Metric IDs whose <code>compliance_relevant: true</code> flag contributed Signal 1
<code>triggered_by_frameworks</code>	Regulatory framework tags from the triggered metrics — e.g. <code>["mifid2"]</code> , <code>["basel3"]</code>
<code>regulatory_trace_id</code>	The trace record identifier written to the ALS regulatory partition for the triggered framework(s)
<code>artifact_set_version</code>	Artifact schema version — used by regulatory consumers to validate the structure
<code>export_requires_lineage</code>	<code>true</code> — consumer must not present export until sealing is confirmed

Field	Description
<code>classification_ceiling_applied</code>	<code>true</code> if a Basel III stress scenario metric triggered the RESTRICTED classification ceiling

Example

The portfolio manager query carries no compliance-relevant metrics — the PAS is not invoked and no compliance block appears in the response.

For a regulatory query against `lcr` and `nsfr` (both `compliance_relevant: true`, `regulatory_framework: ["base13"]`) where the SVL's compliance intent score is `0.94`, both signals are active. The PAS assembles the Provenance Artifact, seals it in the ALS, and returns:

```
{
  "compliance_purpose":      true,
  "intent_score":           0.94,
  "triggered_by_metrics":   ["lcr", "nsfr"],
  "triggered_by_frameworks": ["base13"],
  "regulatory_trace_id":    "trace-20260518-093247-lcr-b3",
  "artifact_set_version":   "1.0",
  "export_requires_lineage": true,
  "classification_ceiling_applied": true
}
```

This block is included in the MCP tool response under the `compliance` key. The consumer renders an appropriate disclosure to the user and withholds export affordances until the platform signals sealing is complete.

Analytical Output Format

The Analytics Engine is headless: it produces no rendered output. Every successful analytical request returns a structured MCP tool response. When narrative synthesis is enabled (`features.narrativeSynthesis: true`), the response contains four elements; when disabled, three.

Output Elements

Element	Field	Always present	Description
Display specification	<code>display_spec</code>	Yes	A Data Visualization Language (DVL) JSON object — either a chart or table specification. Consumers render from this.
Structured result	<code>data</code>	Yes	The computed rows and schema — metric values, dimension values, and units.

Element	Field	Always present	Description
Governed narrative	<code>narrative</code>	No (feature-flag controlled)	A governed summary produced by the NSA — <code>lead</code> (one sentence), <code>detail</code> (2–4 sentences), <code>anchoredTo</code> (result row references). Present when <code>features.narrativeSynthesis</code> is enabled.
Lineage reference	<code>result_id</code> + <code>lineage_url</code>	Yes	A unique result identifier and the URL of the full lineage record
Compliance artifacts	<code>compliance</code>	No (compliance-tier only)	Present when both <code>compliance_relevant</code> metrics are queried AND the SVL classifies query intent as compliance-purpose. Contains regulatory trace ID, triggered metrics, triggered frameworks, intent classification score, <code>export_requires_lineage</code> flag (signals export is blocked until artifact is sealed), and <code>classification_ceiling_applied</code> flag (set when a <code>basel3</code> stress scenario metric triggered RESTRICTED classification ceiling).

Full MCP Response Structure

```
{
  "result_id": "res_20260514_093247_a1b2c3",
  "lineage_url": "https://api.analytics-platform.io/v1/lineage/res_20260514_093247_a1b2c3",
  "display_spec": {
    "type": "chart",
    "mark": "bar",
    ...
  },
  "data": {
    "schema": [ ... ],
    "rows": [ ... ]
  },
  "narrative": {
    "lead": "2 of your 4 equity portfolios outperformed their benchmark this quarter.",
    "detail": "Global Equity Opportunities returned 4.21% against a benchmark of 3.85%. UK Core Income returned 2.87% against its benchmark of 2.54%. Asia Pacific Growth and EUR Balanced Income underperformed, returning 3.67% and 1.93% respectively against benchmarks of 3.90% and 2.31%.",
    "anchoredTo": ["GLOB_EQ_OPP", "UK_CORE_INC", "ASIA_PAC_GRW", "EUR_BAL_INC"]
  },
  "compliance": {
    "compliance_purpose": true,
    "intent_score": 0.94,
    "triggered_by_metrics": ["lcr_ratio", "nsfr_ratio"],
    "triggered_by_frameworks": ["basel3"],
    "regulatory_trace_id": "trace_20260518_093247_lcr",
    "artifact_set_version": "1.0",
    "export_requires_lineage": true,
    "classification_ceiling_applied": true
  },
  "meta": {
    "latencyMs": 1285,
    "cacheHit": false,
    "rowCount": 4,
    "backendsUsed": ["primary-warehouse"],
    "performanceImpactUnits": 620
  }
}
```

The `compliance` block is absent for standard queries. Its presence indicates the full Provenance Artifact was active for this response.

Data Visualization Language (DVL)

DVL is the JSON specification language used for the `display_spec` field. Two types share a consistent discriminated envelope.

Chart specification (`type: "chart"`) — produced when a chart contract matches:

```
{
  "type": "chart",
  "mark": "bar",
  "data": { "values": [ { "portfolio": "Global Equity", "tracking_error": 0.042 }, ... ] },
  "encoding": {
    "x": { "field": "portfolio", "type": "nominal", "title": "Portfolio" },
    "y": { "field": "tracking_error", "type": "quantitative", "title": "Tracking Error" }
  },
  "colorScheme": ["#003f5c", "#58508d", "#bc5090"],
  "formatHints": {
    "tracking_error": { "format": ".2%", "unit": "%" }
  }
}
```

Table specification (`type: "table"`) — produced when no chart contract matches:

```
{
  "type": "table",
  "columns": [
    { "field": "portfolio", "label": "Portfolio", "type": "string" },
    { "field": "tracking_error", "label": "Tracking Error", "type": "number", "format": ".2%" },
    { "field": "limit", "label": "Limit", "type": "number", "format": ".2%" },
    { "field": "breached", "label": "Breached", "type": "boolean" }
  ],
  "data": [ ... ],
  "thresholds": [
    { "field": "tracking_error", "operator": "gt", "reference": "limit", "severity": "warning" }
  ]
}
```

Column labels in table specifications come from SMR metric and dimension `display.label` values — never from physical field names.

Example

The MCP Capability Layer assembles the DVL display specification and the NSA narrative into the final tool response:

```
{
  "jsonrpc": "2.0",
  "id": "req-8a3f2c",
  "result": {
    "content": [
      {
        "type": "text",
        "text": "Across your 4 equity portfolios this quarter, 2 outperformed their benchmark. Global Equity Opportunities led at 4.21% vs 3.85% (+36bps). UK Core Income also outperformed at 2.87% vs 2.54% (+33bps). Asia Pacific Growth and European Balanced Income underperformed, with Asia Pacific Growth the furthest behind at -23bps."
      },
      {
        "type": "resource",
        "resource": {
          "uri": "analytics://result/res-20260518-093247-wk4n",
          "mimeType": "application/vnd.analytics.dvl+json",
          "text": "{ ... DVL grouped bar chart specification ... }"
        }
      }
    ],
    "result_id": "res-20260518-093247-wk4n",
    "lineage_url": "https://api.analytics-platform.io/v1/lineage/res-20260518-093247-wk4n",
    "meta": {
      "latencyMs": 1243,
      "cacheHit": false,
      "rowCount": 4,
      "backendsUsed": ["primary-warehouse"],
      "performanceImpactUnits": 620
    }
  }
}
```

The chat engine renders the grouped bar chart inline and displays the narrative as the assistant's reply. The `result_id` is retained for any follow-up drilldown.

External Components

The following components appear in the architecture diagram and interact with the Analytics Engine but sit outside its boundary. They are not built or owned by the Analytics Engine; they are pre-existing or independently deployed services that the platform integrates with.

Conversational AI — Chat Front End

The AI Chat Platform is the conversational consumer of the Analytics Engine. It relays natural language questions from users to the Analytics Engine and renders the structured results it receives. Intent resolution — identifying which governed operation matches the user's question and binding its parameters — is performed inside the Analytics Engine by the Intent Resolution Agent. The AI Chat Platform performs no NL translation and has no dependency on the SMR operation catalogue.

Responsibilities:

Responsibility	Detail
Query relay	Forwards the user's natural language query and JWT to the Analytics Engine via <code>run_analytics</code> . The query is sent as-is — no structured <code>operation_id</code> or <code>params</code> is constructed by the consumer.
JWT forwarding	Passes the host-issued JWT unmodified with every <code>run_analytics</code> call. The Analytics Engine performs its own JWT validation; the AI Chat Platform does not pre-authorise or modify token claims.
Confirmation card handling	If the Analytics Engine returns a confirmation card (<code>requiresIntentConfirmation: true</code>), the AI Chat Platform renders it to the user and re-submits with <code>confirmed: true</code> when the user approves.
Result rendering	Renders the DVL <code>display_spec</code> returned by the Analytics Engine. Surfaces the governed <code>narrative</code> as the assistant's reply. Retains the <code>result_id</code> for follow-up <code>drilldown</code> calls.

The AI Chat Platform has no access to physical schemas, execution backends, or metric definitions. Entitlement enforcement, intent resolution, query planning, and execution are entirely the Analytics Engine's responsibility.

Semantic Data Repository (SDR)

The Semantic Data Repository is a pre-existing organisational component — the governed store of JSON-based data metadata definitions that describe the organisation's information assets. It exists independently of the Analytics Platform and is not built or owned by it. For most organisations it will already exist before the Analytics Platform is deployed.

The SDR contains the organisation's foundational data context: data models, object models, critical data elements, quality rules, physical schemas, and data lineage records — *what data exists and how it is structured*. The SMR is a separate store for metric metadata definitions — *what the data means analytically* and how it should be calculated, aggregated, and governed. Both are independent stores housed within the Data Context Store (DCS), which is the external container for all governed context.

Role in the platform:

Function	Detail
Foundational data context	Provides the data model and schema metadata definitions that SMR metric <code>physical_mapping</code> fields reference.
Physical mapping source	The <code>physical_mapping</code> fields in SMR metric definitions resolve against SDR schema metadata to identify the physical tables and columns that back each metric.
Search and RAG	The DCS search index (spanning both SMR and SDR) supports the <code>list_operations</code> tool and the IRA's vector similarity search over SMR operation and metric embeddings for NL intent resolution.

Data Entitlements Store (DES)

The Data Entitlements Store (DES) is an independent external component that holds the organisation's data and analytics entitlement policies. It is managed separately from the Analytics Engine and from the data platform — it is a dedicated governance control point.

Entitlement policies are declared at the **logical object and data element level**. A policy grants or restricts access to a named metric, a named dimension, or a named data element as governed concepts. Policies do not reference physical tables, schemas, column names, or connection strings. This separation means entitlements remain stable as the underlying physical implementation evolves, and they are comprehensible to business data owners, compliance teams, and governance teams who have no visibility into the data platform's internal structure.

Role in the platform:

Function	Detail
Role definition storage	Stores <code>analytical_role</code> definitions: metric access sets, dimension access sets, row predicate templates, column masks, and classification ceilings per role.
Logical-level policy	All policies reference SMR-registered metric and dimension identifiers — never physical table or column names.
RAPL lookup	RAPL reads role definitions from the Data Entitlements Store (DES) at query time, keyed on the role claims extracted from the caller's JWT.
Independent governance	Managed by the Entitlements Manager persona. Changes to entitlement policies do not require changes to metric definitions, platform configuration, or backend schemas.

vega2img

`vega2img` is an optional, independently deployed MCP render service. It converts the Analytics Engine's DVL display specification into a static image — SVG or PNG — for consumers that cannot natively render a DVL specification inline.

Integration model:

Property	Detail
Deployment	Standalone service, deployed and registered independently of the Analytics Engine.
Registration	Registered directly with the AI consumer (AI Chat Platform, agent, or custom application) as a peer MCP server — not as part of the Analytics Engine.
Invocation	The consumer calls <code>vega2img</code> as a separate tool invocation, passing the <code>display_spec</code> returned by the Analytics Engine's <code>run_analytics</code> response.
Scope	Stateless render only. <code>vega2img</code> has no access to the Analytics Engine, the SMR, or any execution backend. It receives a self-contained DVL specification and returns an image.

`vega2img` is not required for consumers that can natively render DVL specifications (most AI chat platforms and modern front-end frameworks). It is included in the architecture diagram for completeness — agentic pipelines that produce static report output are the primary use case.